

Lab 4: Data Visualization and EDA

CPE232 Data Models

1. Load all Superstore datasets.

Note: The same datasets used in Lab 3

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
import plotly.io as pio
pio.renderers.default = "png+notebook_connected+vscode"
```

```
In [2]: # Write your code here
superstore_order = pd.read_csv('Superstore/superstore_order.csv')
superstore_people = pd.read_csv('Superstore/superstore_people.csv')
superstore_return = pd.read_csv('Superstore/superstore_return.csv')
```

2. Determine shape of each dataset (print out the results as well).

```
In [3]: # Write your code here
print(superstore_order.shape)
print(superstore_people.shape)
print(superstore_return.shape)

(8880, 21)
(4, 2)
(296, 2)
```

3. Show information of the dataset.

```
In [4]: # Write your code here
superstore_order.info()

<class 'pandas.DataFrame'>
RangeIndex: 8880 entries, 0 to 8879
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Row ID           8880 non-null  int64
1   Order ID         8880 non-null  str
2   Order Date       8880 non-null  str
3   Ship Date        8880 non-null  str
4   Ship Mode        8880 non-null  str
5   Customer ID      8880 non-null  str
6   Customer Name    8880 non-null  str
7   Segment         8880 non-null  str
8   Country          8880 non-null  str
9   City            8880 non-null  str
10  State            8880 non-null  str
11  Postal Code      8880 non-null  int64
12  Region           8880 non-null  str
13  Product ID       8880 non-null  str
14  Category         8880 non-null  str
15  Sub-Category     8880 non-null  str
16  Product Name     8880 non-null  str
17  Sales            8880 non-null  float64
18  Quantity         8880 non-null  int64
19  Discount         8880 non-null  float64
20  Profit           8880 non-null  float64
dtypes: float64(3), int64(3), str(15)
memory usage: 1.4 MB
```

```
In [5]: superstore_people.info()

<class 'pandas.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0   Person  4 non-null      str
1   Region  4 non-null      str
dtypes: str(2)
memory usage: 196.0 bytes
```

```
In [6]: superstore_return.info()

<class 'pandas.DataFrame'>
RangeIndex: 296 entries, 0 to 295
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0   Returned 296 non-null   str
1   Order ID 296 non-null   str
dtypes: str(2)
memory usage: 4.8 KB
```

4. Are there any missing values? If so, in which column?

Ans: There are no missing values in any column

- 5.

- 5.1 List unique segments
- 5.2 List unique segments and their corresponding count
- 5.3 Create a pie chart to demonstrate unique segments and their count
- 5.4 Briefly describe what could be interpreted from this pie chart

Note: please create additional cells to answer 5.2 - 5.3

```

In [7]: # Write your code here (5.1)
uq_segment = superstore_order["Segment"].unique()
uq_segment

Out[7]: <StringArray>
['Consumer', 'Corporate', 'Home Office']
Length: 3, dtype: str


In [8]: superstore_order["Segment"].value_counts()

Out[8]: Segment
Consumer      4613
Corporate     2673
Home Office   1594
Name: count, dtype: int64

In [9]: fig = px.pie(superstore_order, names='Segment', title='Super Store Order Segment Proportion', hole=0.3)
fig.show()

```

Super Store Order Segment Proportion



Answer for the question 5.4

Ans: The Consumer segment is the largest and accounts for more than half of the total orders.

6.

- 6.1 List unique states
- 6.2 List top-10 unique states and their corresponding count
- 6.3 Create a bar chart (vertical) to demonstrate the count of top-10 unique states
- 6.4 Based on 6.2, also include the total sales of these states (show your result as a dataframe)
- 6.5 Using the result from 6.4, if you were the owner of this superstore, what information could be interpreted from this result?

Note: please create additional cells to answer 6.2 - 6.4

```

In [10]: # Write your code here (6.1)
superstore_order["State"].unique()

Out[10]: <StringArray>
[ 'Kentucky',      'California',      'Florida',
  'North Carolina', 'Washington',      'Texas',
  'Wisconsin',      'Utah',            'Nebraska',
  'Pennsylvania',   'Illinois',        'Minnesota',
  'Michigan',       'Delaware',        'Indiana',
  'New York',       'Arizona',         'Virginia',
  'Tennessee',      'Alabama',         'South Carolina',
  'Oregon',         'Colorado',        'Iowa',
  'Ohio',           'Missouri',        'Oklahoma',
  'New Mexico',     'Louisiana',       'Connecticut',
  'New Jersey',     'Massachusetts',   'Georgia',
  'Nevada',         'Rhode Island',    'Mississippi',
  'Arkansas',       'Montana',         'New Hampshire',
  'Maryland',       'District of Columbia', 'Kansas',
  'Vermont',        'Maine',           'South Dakota',
  'Idaho',          'North Dakota',    'Wyoming',
  'West Virginia']
Length: 49, dtype: str

In [11]: top_state = superstore_order["State"].value_counts().head(10).reset_index()
top_state


```

```
Out[11]:
```

	State	count
0	California	1754
1	New York	1001
2	Texas	860
3	Pennsylvania	531
4	Washington	452
5	Illinois	427
6	Ohio	396
7	Florida	339
8	Michigan	230
9	North Carolina	229

```
In [12]: px.bar(top_state,x = "State", y = "count", title="Order by State")
```

Order by State



Answer for the question 6.5

Ans: There are only 4 State that have orders more than 500 orders.

7.

- 7.1 List unique categories
- 7.2 Create a bar chart (horizontal) to demonstrate the proportion of these categories
- 7.3 Compute the ratio of these categories in percentage and print the results


Note: please create additional cells to answer 7.2 - 7.3

```
In [13]: # Write your code here (7.1)
superstore_order["Category"].unique()
```

```
Out[13]: <StringArray>
['Furniture', 'Office Supplies', 'Technology']
Length: 3, dtype: str
```

```
In [14]: cat_count = superstore_order["Category"].value_counts().reset_index()
fig = px.bar(cat_count,x="count",y = "Category",title = "Order by Category",orientation="h")
fig.show()
```

Order by Category



```
In [15]: order_count = cat_count["count"].sum()
Technology_ratio = cat_count[cat_count["Category"] == "Technology"]["count"].item() / order_count
Furniture_ratio = cat_count[cat_count["Category"] == "Furniture"]["count"].item() / order_count
Office_Supplies_ratio = cat_count[cat_count["Category"] == "Office Supplies"]["count"].item() / order_count
print(f"Technology_ratio {Technology_ratio:.2%}")
print(f"Furniture_ratio {Furniture_ratio:.2%}")
print(f"Office_Supplies_ratio {Office_Supplies_ratio:.2%}")

Technology_ratio 18.47%
Furniture_ratio 21.17%
Office_Supplies_ratio 60.36%
```

8. Update the type of columns that contain dates to *datetime* and show information after an update.

```
In [16]: # write your code here
superstore_order["Order Date"] = pd.to_datetime(superstore_order["Order Date"], format="%d/%m/%Y")
superstore_order["Ship Date"] = pd.to_datetime(superstore_order["Ship Date"], format="%d/%m/%Y")
superstore_order.info()

<class 'pandas.DataFrame'>
RangeIndex: 8880 entries, 0 to 8879
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Row ID           8880 non-null   int64
1   Order ID         8880 non-null   str
2   Order Date       8880 non-null   datetime64[us]
3   Ship Date        8880 non-null   datetime64[us]
4   Ship Mode        8880 non-null   str
5   Customer ID      8880 non-null   str
6   Customer Name    8880 non-null   str
7   Segment          8880 non-null   str
8   Country          8880 non-null   str
9   City            8880 non-null   str
10  State            8880 non-null   str
11  Postal Code      8880 non-null   int64
12  Region           8880 non-null   str
13  Product ID       8880 non-null   str
14  Category         8880 non-null   str
15  Sub-Category     8880 non-null   str
16  Product Name     8880 non-null   str
17  Sales            8880 non-null   float64
18  Quantity         8880 non-null   int64
19  Discount         8880 non-null   float64
20  Profit          8880 non-null   float64
dtypes: datetime64[us](2), float64(3), int64(3), str(13)
memory usage: 1.4 MB
```

9. Create a new column "Processing time day" to show number of days taken to ship an order and show your result in a dataframe format.

Hint: The duration starts as soon as the item has been ordered and ends once the order has successfully shipped.

```
In [17]: # write your code here
superstore_order["Processing time day"] = superstore_order["Ship Date"] - superstore_order["Order Date"]
superstore_order["Processing time day"]
```

```
Out[17]: 0      3 days
1      3 days
2      4 days
3      7 days
4      7 days
...
8875   6 days
8876   6 days
8877   6 days
8878   6 days
8879   2 days
Name: Processing time day, Length: 8880, dtype: timedelta64[us]
```

10. Based on the result in 9.

- 10.1 How many orders are there that take more than 5 days to process?
- 10.2 Show the top 5 rows (expected output should contain these columns: Order ID, Order Date, Ship Date, Processing time day, Quantity)
- 10.3 Plot the histogram based on the column Quantity

Note: please create additional cells to answer 10.2 - 10.3

```
In [18]: # Write your code here (10.1)
(superstore_order["Processing time day"].dt.days > 5).sum()
```


```
Out[18]: np.int64(1656)
```

```
In [19]: superstore_order.loc[:,4,["Order ID","Order Date","Ship Date","Processing time day","Quantity"]]
```

```
Out[19]:
```

	Order ID	Order Date	Ship Date	Processing time day	Quantity
0	CA-2016-152156	2016-11-08	2016-11-11	3 days	2
1	CA-2016-152156	2016-11-08	2016-11-11	3 days	3
2	CA-2016-138688	2016-06-12	2016-06-16	4 days	2
3	US-2015-108966	2015-10-11	2015-10-18	7 days	5
4	US-2015-108966	2015-10-11	2015-10-18	7 days	2


```
In [20]: plt.figure(figsize=(10, 6))
plt.hist(superstore_order['Quantity'], bins=10, color='yellow', edgecolor='black')
plt.title("Order Quantity Histogram")
plt.xlabel('Quantity')
plt.ylabel('Frequency')
plt.show()
```



11. Total sales compare across different regions

- 11.1 Create a bar chart to visualize.

```
In [21]: # Write your code here (11.1)
reg_count = superstore_order["Region"].value_counts().reset_index()
px.bar(reg_count, x = "Region", y = "count", title="Order by Region")
```



- 11.2 How do total sales compare across different regions? Explain in as much detail as possible.


Ans: South region is only region that order less than 2000 orders while West and East have order more than 2500 orders.

12. Which states have the highest number of returns? Use a horizontal bar chart.

Ans: California


```
In [22]: # Write your code here (12)
returned_order = pd.concat([superstore_order, superstore_return], axis=1, join='inner')
returned_order_count_by_state = returned_order["State"].value_counts().reset_index()
px.bar(returned_order_count_by_state, x = "State", y = "count", title="Returned Order by State")
```

Returned Order by State



13. What is the correlation between numerical variables in the superstore_order dataset? Use a heatmap *Hint: Use seaborn to create a heatmap.*

```
In [23]: # Write your code here (13)
numeric_superstore = superstore_order.select_dtypes(include=['number'])
correlation_matrix = numeric_superstore.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Correlation Heatmap of Superstore Order Variables')
plt.show()
```



14. Create a USA State-Level Choropleth Map to visualize total sales per state.

- The darkest color represents the highest total sales.
- The lightest color represents the lowest total sales.
- Use a continuous gradient scale (e.g., dark blue to light blue, dark red to light red, or any custom gradient of your choice).

Hint: Use plotly.express


```
In [24]: # Write your code here (14)
state_sales = superstore_order.groupby("State")["Sales"].sum().reset_index()

state_map = {
    'Alabama': 'AL', 'Alaska': 'AK', 'Arizona': 'AZ', 'Arkansas': 'AR', 'California': 'CA',
    'Colorado': 'CO', 'Connecticut': 'CT', 'Delaware': 'DE', 'Florida': 'FL', 'Georgia': 'GA',
    'Hawaii': 'HI', 'Idaho': 'ID', 'Illinois': 'IL', 'Indiana': 'IN', 'Iowa': 'IA',
    'Kansas': 'KS', 'Kentucky': 'KY', 'Louisiana': 'LA', 'Maine': 'ME', 'Maryland': 'MD',
    'Massachusetts': 'MA', 'Michigan': 'MI', 'Minnesota': 'MN', 'Mississippi': 'MS',
    'Missouri': 'MO', 'Montana': 'MT', 'Nebraska': 'NE', 'Nevada': 'NV', 'New Hampshire': 'NH',
    'New Jersey': 'NJ', 'New Mexico': 'NM', 'New York': 'NY', 'North Carolina': 'NC',
    'North Dakota': 'ND', 'Ohio': 'OH', 'Oklahoma': 'OK', 'Oregon': 'OR', 'Pennsylvania': 'PA',
    'Rhode Island': 'RI', 'South Carolina': 'SC', 'South Dakota': 'SD', 'Tennessee': 'TN',
    'Texas': 'TX', 'Utah': 'UT', 'Vermont': 'VT', 'Virginia': 'VA', 'Washington': 'WA',
    'West Virginia': 'WV', 'Wisconsin': 'WI', 'Wyoming': 'WY', 'District of Columbia': 'DC'
}

state_sales['State Code'] = state_sales['State'].map(state_map)
```

```
fig = px.choropleth(
    state_sales,
    locations='State Code',
    locationmode='USA-states',
    color='Sales',
    hover_name='State',
    color_continuous_scale='Blues',
    scope='usa',
    title='U.S. State Sales'
)
fig.show()
```

U.S. State Sales



14.2 Answer the following questions:

1. Which state has the highest total sales?
2. How do sales anomalies affect the gradient color shading on the map?
3. If you change the color scale, does it impact readability? Why or why not?

Ans:

1. California
2. gradient Scale of anomalies are compute as darkest or lightest while other is on similar gradient lead to challenges for analyzing.
3. Changing color sclae impact readability, a sequential scale is ideal for showing ordered, low-to-high data. A diverging scale is better for highlighting extremes or a significant value.


15. Create a box plot to compare the different shipping modes based on total profit.

```
In [25]: #Write your code here (15)
superstore_order["Ship Mode"].unique()
```

```
Out[25]: <StringArray>
['Second Class', 'Standard Class', 'First Class', 'Same Day']
Length: 4, dtype: str
```

```
In [33]: fig = px.box(superstore_order,x="Ship Mode",y="Profit",title="Profit distribution",log_y=True)
fig.show()
```

Profit distribution



15.2 Which shipping mode has the highest median profit?

Ans: Second Class

[BONUS 20 pts] Determine the percentage of customers who:

- B1) returned the product once
- B2) returned the product at least once

- B3) never returned the product
- Finally, Plot a comparison of B2 and B3

Note: please create additional cells to answer the above points

```
In [27]: # B1
customer_count = len(superstore_order["Customer ID"].unique())
concated_return = pd.concat([superstore_order, superstore_return], axis=1, join='inner')
return_count = concat_return["Customer ID"].value_counts()
once_return = (return_count == 1).sum()
B1 = once_return/customer_count
print(f"B1 = {B1:.2%}")
```

B1 = 7.10%

```
In [28]: # B2
return_count = concat_return["Customer ID"].value_counts()
B2 = (return_count.sum())/customer_count
print(f"B2 = {B2:.2%}")
```

B2 = 37.52%

```
In [29]: # B3
B3 = (customer_count - return_count.sum())/customer_count
print(f"B3 = {B3:.2%}")
```

B3 = 62.48%

```
In [30]: labels = ['B2', 'B3']
percentage = [B2, B3]
plt.bar(labels, percentage)
plt.title('Comparison of B2 and B3')
plt.xlabel('Variable')
plt.ylabel('Percentage')
plt.show()
```

