

Lab 3: Data Preparation

CPE232 Data Models

For this lab, we will learn how to perform data preparation for data analysis after receiving the raw data from the data source.



[1] Reviews on Pandas

1.1) Discover

- methods to explore and understand your DataFrame

```
In [72]: import pandas as pd
```

```
df = pd.read_csv('nss15.csv')
df.head()
```

Out[72]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	body
0	150733174	7/11/2015	15.7762	V	5	Male	NaN		57
1	150734723	7/6/2015	83.2157	S	36	Male	White		57
2	150817487	8/2/2015	74.8813	L	20	Female	NaN		71
3	150717776	6/26/2015	15.7762	V	61	Male	NaN		71
4	150721694	7/4/2015	74.8813	L	88	Female	Other		62

In [73]:

```
# see the shape of the dataframe
print(df.shape)
```

(334839, 12)

In [74]:

```
# seeing the summary of the dataframe
print(df.info())
```

```
<class 'pandas.DataFrame'>
RangeIndex: 334839 entries, 0 to 334838
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   caseNumber      334839 non-null   int64  
 1   treatmentDate   334839 non-null   str    
 2   statWeight      334839 non-null   float64 
 3   stratum         334839 non-null   str    
 4   age              334839 non-null   int64  
 5   sex              334837 non-null   str    
 6   race             205014 non-null   str    
 7   diagnosis        334839 non-null   int64  
 8   bodyPart         334839 non-null   int64  
 9   disposition      334839 non-null   int64  
 10  location         334839 non-null   int64  
 11  product          334839 non-null   int64  
dtypes: float64(1), int64(7), str(4)
memory usage: 30.7 MB
None
```

In [75]:

```
# seeing the stats of the column in dataframe
df.describe()
```

Out[75]:

	caseNumber	statWeight	age	diagnosis	bodyPart	dispo
count	3.348390e+05	334839.000000	334839.000000	334839.000000	334839.000000	334839.000000
mean	1.510271e+08	39.343028	31.385451	60.154591	64.374192	1.3
std	1.720330e+06	34.142933	26.105098	6.170699	24.002331	0.9
min	1.501032e+08	4.965500	0.000000	41.000000	0.000000	1.0
25%	1.504405e+08	15.059100	10.000000	57.000000	35.000000	1.0
50%	1.507358e+08	15.776200	23.000000	59.000000	75.000000	1.0
75%	1.510231e+08	74.881300	51.000000	64.000000	82.000000	1.0
max	1.603418e+08	97.923900	107.000000	74.000000	94.000000	9.0



1.2) Selecting variables

- select specific columns from the DataFrame to create a new DataFrame with only those columns

In [76]: `# select columns based on the data type
df.select_dtypes(include=['number'])`

Out[76]:

	caseNumber	statWeight	age	diagnosis	bodyPart	disposition	location	product
0	150733174	15.7762	5	57	33	1	9	1267
1	150734723	83.2157	36	57	34	1	1	1439
2	150817487	74.8813	20	71	94	1	0	3274
3	150717776	15.7762	61	71	35	1	0	611
4	150721694	74.8813	88	62	75	1	0	1893
...
334834	150739278	15.0591	7	59	76	1	1	1864
334835	150733393	5.6748	3	68	85	1	0	1931
334836	150819286	15.7762	38	71	79	1	0	3250
334837	150823002	97.9239	38	59	82	1	1	464
334838	150723074	49.2646	5	57	34	1	9	3273

334839 rows × 8 columns



In [77]: `# select column by .loc
df.loc[:, 'treatmentDate':'diagnosis']`

Out[77]:

	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	7/11/2015	15.7762	V	5	Male	NaN	57
1	7/6/2015	83.2157	S	36	Male	White	57
2	8/2/2015	74.8813	L	20	Female	NaN	71
3	6/26/2015	15.7762	V	61	Male	NaN	71
4	7/4/2015	74.8813	L	88	Female	Other	62
5	7/2/2015	5.6748	C	1	Female	White	71
6	6/8/2015	15.7762	V	25	Male	Black	51

In [78]: `# select row by .iloc
df.iloc[0:5]`

Out[78]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	body
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	

1.3) Filtering the data

In [79]: `# filter rows based on the condition
df[df['age'] > 50]`

Out[79]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62
7	150704114	6/14/2015	83.2157	S	53	Male	White	57
8	150736558	7/16/2015	83.2157	S	98	Male	Black	59
16	150901411	8/27/2015	83.2157	S	65	Female	White	59
...
334811	150702215	6/27/2015	15.7762	V	51	Female	NaN	53
334815	151100368	10/28/2015	83.2157	S	85	Female	NaN	57
334819	150528367	1/13/2015	49.2646	M	85	Female	NaN	57
334826	150648619	6/17/2015	15.7762	V	52	Female	White	64
334829	150633526	4/4/2015	49.2646	M	51	Female	NaN	56

85235 rows × 12 columns



In [80]: `# filter coloum based on column name
df.filter(like='age')`

Out[80]:

	age
0	5
1	36
2	20
3	61
4	88
...	...
334834	7
334835	3
334836	38
334837	38
334838	5

334839 rows × 1 columns

1.4) Sorting

- Sort the DataFrame by its index based on column

```
In [81]: # sort the dataframe based on column name and ascending order
df.sort_values(by='statWeight', ascending=False)
```

Out[81]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
275174	150343700	3/9/2015	97.9239	M	48	Male	NaN	57
36	151029422	10/6/2015	97.9239	M	37	Male	White	64
334806	150612491	5/29/2015	97.9239	M	18	Female	White	59
334810	150725804	7/8/2015	97.9239	M	33	Female	Black	71
275161	150450816	4/13/2015	97.9239	M	24	Male	White	71
...
44011	160222258	12/29/2015	4.9655	C	2	Female	Other	71
325320	151213065	11/29/2015	4.9655	C	16	Female	White	62
43891	160113865	12/28/2015	4.9655	C	4	Male	White	59
43628	151130111	11/9/2015	4.9655	C	13	Male	Black	53
43523	151139237	11/16/2015	4.9655	C	2	Female	Black	57

334839 rows × 12 columns



```
In [82]: # sort the index of the dataframe
df.sort_index()
```

Out[82]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57
1	150734723	7/6/2015	83.2157	S	36	Male	White	57
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57

334839 rows × 12 columns



1.5) Add/Remove

- This section shows how to manipulate the DataFrame's structure

In [83]:

```
# Dropping the column
df.drop(columns=['disposition'])
```

Out[83]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57
1	150734723	7/6/2015	83.2157	S	36	Male	White	57
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57

334839 rows × 11 columns



In [84]:

```
# Adding column and create into a new column
df.assign(new_column=df['diagnosis'] + df['bodyPart'])
```

Out[84]:

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57
1	150734723	7/6/2015	83.2157	S	36	Male	White	57
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57

334839 rows × 13 columns



```
In [85]: # Removing the column and assigning it to a new variable  
ages = df.pop('age')
```

1.6) Clean missing

- to remove rows with missing values or replace missing values with a specified value

```
In [86]: # replacing the missing values with a specified value  
df.fillna(value=0)
```

```
Out[86]:
```

	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	body
0	150733174	7/11/2015	15.7762	V	Male	0	57	
1	150734723	7/6/2015	83.2157	S	Male	White	57	
2	150817487	8/2/2015	74.8813	L	Female	0	71	
3	150717776	6/26/2015	15.7762	V	Male	0	71	
4	150721694	7/4/2015	74.8813	L	Female	Other	62	
...
334834	150739278	5/31/2015	15.0591	V	Male	0	59	
334835	150733393	7/11/2015	5.6748	C	Female	Black	68	
334836	150819286	7/24/2015	15.7762	V	Male	0	71	
334837	150823002	8/8/2015	97.9239	M	Female	White	59	
334838	150723074	6/20/2015	49.2646	M	Female	White	57	

334839 rows × 11 columns



```
In [87]: # Remove the rows with missing values  
df.dropna()
```

Out[87]:

	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	body
1	150734723	7/6/2015	83.2157	S	Male	White		57
4	150721694	7/4/2015	74.8813	L	Female	Other		62
5	150721815	7/2/2015	5.6748	C	Female	White		71
6	150713483	6/8/2015	15.7762	V	Male	Black		51
7	150704114	6/14/2015	83.2157	S	Male	White		57
...
334830	150628863	6/8/2015	15.7762	V	Female	White		64
334831	150607637	5/22/2015	5.6748	C	Female	Black		59
334835	150733393	7/11/2015	5.6748	C	Female	Black		68
334837	150823002	8/8/2015	97.9239	M	Female	White		59
334838	150723074	6/20/2015	49.2646	M	Female	White		57

205014 rows × 11 columns



[2] Data Cleaning and Preparation

.isnull, .dropna, .fillna

2.1) checking

In [88]: `df.columns`

Out[88]: `Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'sex', 'race', 'diagnosis', 'bodyPart', 'disposition', 'location', 'product'], dtype='str')`

In [89]: `# isnull checking
df.isnull().sum()`

```
Out[89]: caseNumber      0
treatmentDate      0
statWeight        0
stratum          0
sex              2
race            129825
diagnosis        0
bodyPart         0
disposition       0
location         0
product          0
dtype: int64
```

```
In [90]: # percentage of missing values for the race
df.race.isnull().sum()/df.shape[0]*100
```

```
Out[90]: np.float64(38.772365226272925)
```

```
In [91]: df.shape[0]
```

```
Out[91]: 334839
```

2.2) Drop column

```
In [92]: # remove column by using
df = df.drop(columns=['race'])
```

```
In [93]: df.head()
```

```
Out[93]:   caseNumber  treatmentDate  statWeight  stratum    sex  diagnosis  bodyPart  disposition
0     150733174    7/11/2015    15.7762      V  Male      57      33
1     150734723    7/6/2015    83.2157      S  Male      57      34
2     150817487    8/2/2015    74.8813      L Female     71      94
3     150717776    6/26/2015    15.7762      V  Male      71      35
4     150721694    7/4/2015    74.8813      L Female     62      75
```



2.3) Data imputation

```
In [94]: # fillna
df['age'] = df['age'].fillna(df['age'].median())
```

```

-----
KeyError                                         Traceback (most recent call last)
File c:\Users\win25\Desktop\Desktop\work\CPE232 Data models\.venv\Lib\site-packages
\pandas\core\indexes\base.py:3641, in Index.get_loc(self, key)
    3640     try:
-> 3641         return self._engine.get_loc(casted_key)
    3642     except KeyError as err:
    3643
    3644     if isna(key):
    3645         return self._get_na_locs(key)

File pandas\_libs/index.pyx:168, in pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/index.pyx:197, in pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/hashtable_class_helper.pxi:7668, in pandas._libs.hashtable.PyObjec
tHashTable.get_item()

File pandas/_libs/hashtable_class_helper.pxi:7676, in pandas._libs.hashtable.PyObjec
tHashTable.get_item()

KeyError: 'age'

The above exception was the direct cause of the following exception:

KeyError                                         Traceback (most recent call last)
Cell In[94], line 2
      1 # fillna
-> 2 df['age'] = df[      ].fillna(df['age'].median())

File c:\Users\win25\Desktop\Desktop\work\CPE232 Data models\.venv\Lib\site-packages
\pandas\core\frame.py:4378, in DataFrame.__getitem__(self, key)
    4376     if self.columns.nlevels > 1:
    4377         return self._getitem_multilevel(key)
-> 4378     indexer = self.columns.get_loc(key)
    4379     if is_integer(indexer):
    4380         indexer = [indexer]

File c:\Users\win25\Desktop\Desktop\work\CPE232 Data models\.venv\Lib\site-packages
\pandas\core\indexes\base.py:3648, in Index.get_loc(self, key)
    3643     if isinstance(casted_key, slice) or (
    3644         isinstance(casted_key, abc.Iterable)
    3645         and any(isinstance(x, slice) for x in casted_key)
    3646     ):
    3647         raise InvalidIndexError(key) from err
-> 3648     raise KeyError(key) from err
    3649 except TypeError:
    3650     # If we have a listlike key, _check_indexing_error will raise
    3651     # InvalidIndexError. Otherwise we fall through and re-raise
    3652     # the TypeError.
    3653     self._check_indexing_error(key)

KeyError: 'age'

```

[Q1] From the above cell, Why it showing an error?

Ans: Because Dataframe df don't have Column name "age" so it will occurs error when trying to access using df["age"]

[Q2] Fix the error from Q1 problem.

```
In [ ]: # [Q2]

# hint: see the cell that run `df.pop()`
df["age"] = ages

# fillna again
df['age'] = df['age'].fillna(df['age'].median())

df.head()
```

```
Out[ ]:   caseNumber  treatmentDate  statWeight  stratum    sex  diagnosis  bodyPart  disposi
0      150733174     7/11/2015    15.7762      V  Male       57        33
1      150734723     7/6/2015     83.2157      S  Male       57        34
2      150817487     8/2/2015    74.8813      L Female      71        94
3      150717776     6/26/2015    15.7762      V  Male       71        35
4      150721694     7/4/2015    74.8813      L Female      62        75
```



2.4) Drop row that have missing value

```
In [ ]: # remove column by using .dropna()
df = df.dropna()
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: caseNumber      0
treatmentDate      0
statWeight         0
stratum            0
sex                0
diagnosis          0
bodyPart           0
disposition         0
location           0
product            0
age                0
dtype: int64
```

Datetime

2.5) Working with the datetime format

```
In [ ]: df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
```

```
In [ ]: df.info()
```

```
<class 'pandas.DataFrame'>
Index: 334837 entries, 0 to 334838
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   caseNumber       334837 non-null   int64  
 1   treatmentDate    334837 non-null   datetime64[us] 
 2   statWeight       334837 non-null   float64 
 3   stratum          334837 non-null   str    
 4   sex              334837 non-null   str    
 5   diagnosis        334837 non-null   int64  
 6   bodyPart         334837 non-null   int64  
 7   disposition      334837 non-null   int64  
 8   location          334837 non-null   int64  
 9   product          334837 non-null   int64  
 10  age              334837 non-null   int64  
dtypes: datetime64[us](1), float64(1), int64(7), str(2)
memory usage: 30.7 MB
```

```
In [ ]: df['Year'] = df['treatmentDate'].dt.year
```

```
In [ ]: df['Month'] = df['treatmentDate'].dt.month
```

```
In [ ]: df.head()
```

```
Out[ ]:   caseNumber  treatmentDate  statWeight  stratum  sex  diagnosis  bodyPart  disposition
0   150733174     2015-07-11    15.7762      V   Male      57      33
1   150734723     2015-07-06    83.2157      S   Male      57      34
2   150817487     2015-08-02    74.8813      L Female     71      94
3   150717776     2015-06-26    15.7762      V   Male      71      35
4   150721694     2015-07-04    74.8813      L Female     62      75
```



[Q3] Can you change the format to DD/MM/YYYY? Show your work.

```
In [ ]: df['treatmentDateNewFormat'] = df['treatmentDate'].dt.strftime('%d/%m/%Y')
```

Out[]:

	caseNumber	treatmentDate	statWeight	stratum	sex	diagnosis	bodyPart	d
0	150733174	2015-07-11	15.7762	V	Male	57	33	
1	150734723	2015-07-06	83.2157	S	Male	57	34	
2	150817487	2015-08-02	74.8813	L	Female	71	94	
3	150717776	2015-06-26	15.7762	V	Male	71	35	
4	150721694	2015-07-04	74.8813	L	Female	62	75	
...
334834	150739278	2015-05-31	15.0591	V	Male	59	76	
334835	150733393	2015-07-11	5.6748	C	Female	68	85	
334836	150819286	2015-07-24	15.7762	V	Male	71	79	
334837	150823002	2015-08-08	97.9239	M	Female	59	82	
334838	150723074	2015-06-20	49.2646	M	Female	57	34	

334837 rows × 14 columns



Combine Dataframe by .merge and .concat

2.6 Merge

In []:

```
import pandas as pd

superstore_order = pd.read_csv('Superstore/superstore_order.csv')
superstore_people = pd.read_csv('Superstore/superstore_people.csv')
superstore_return = pd.read_csv('Superstore/superstore_return.csv')
```

In []:

```
superstore_order.head()
```

Out[]:

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Count
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States
4	5	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States

5 rows × 21 columns



In []: `superstore_people.head()`

Out[]:

	Person	Region
0	Anna Andreadi	West
1	Chuck Magee	East
2	Kelly Williams	Central
3	Cassandra Brandow	South

In []: `superstore_return.head()`

Out[]:

	Returned	Order ID
0	Yes	CA-2017-153822
1	Yes	CA-2017-129707
2	Yes	CA-2014-152345
3	Yes	CA-2015-156440
4	Yes	US-2017-155999

In []:

```
superstore_order.merge(superstore_return[superstore_return["Returned"]=="Yes"],
on="Order ID" ,
how="inner")\
[[ "Customer ID", "Returned"]]\n.drop_duplicates()
```

Out[]:

	Customer ID	Returned
0	ZD-21925	Yes
3	TB-21055	Yes
10	JS-15685	Yes
13	LC-16885	Yes
20	BS-11755	Yes
...
688	ED-13885	Yes
689	TS-21205	Yes
696	MF-17665	Yes
702	SH-19975	Yes
705	RB-19435	Yes

222 rows × 2 columns

[Q4] What does the argument `how="inner"` do?

Ans: ' how = "inner" ' is how to merge 2 table between superstore_order and superstore_return[superstore_return["Returned"]=="Yes"]

"inner" mean these 2 table merge using inner join (CustomerID) after joining we get superstore_return Dataframe with "Returned" Column
from superstore_return[superstore_return["Returned"]=="Yes" Dataframe

[Q5] In your opinion, what information that the result above conveys?

Ans: It shows CustomerID of Customer who returns the order to supermarket

2.7) Concatenate

```
In [ ]: pd.concat([superstore_order, superstore_people], axis=1, join='inner')
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Count
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States

4 rows × 23 columns

```
In [ ]: pd.concat([superstore_order, superstore_people], axis=1, join='outer')
```

Out[]:

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Co
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	l
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	l
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	l
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	l
4	5	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	l
...
8875	8876	US-2016-141264	13/08/2016	19/08/2016	Standard Class	CT-11995	Carol Triggs	Consumer	l
8876	8877	US-2016-141264	13/08/2016	19/08/2016	Standard Class	CT-11995	Carol Triggs	Consumer	l
8877	8878	CA-2017-126928	17/09/2017	23/09/2017	Standard Class	GZ-14470	Gary Zandusky	Consumer	l
8878	8879	CA-2017-126928	17/09/2017	23/09/2017	Standard Class	GZ-14470	Gary Zandusky	Consumer	l

Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Co
8879	8880	US-2015-107944	23/03/2015	25/03/2015	First Class	AM-10360	Alice McCarthy	Corporate

8880 rows × 23 columns

[Q6] What is the difference between `inner` and `outer` on parameter `join` in `pd.concat` ?

Ans 'inner' is inner join between 2 Dataframe of `pd.concat`, inner join works like intersection its join two Dataframe without having NaN value

while 'outer' is outer join between 2 Dataframe of `pd.concat`, outer join works like union after join sometimes new Dataframe will have empty entity

[Q7] Append new record at the end of `superstore_order`. Set the 'Row ID' to your student ID and mock other required fields based on existing entries.

```
In [ ]: # write your code here
from datetime import datetime
mydata = {
    "Row ID" : 67070501021,
    "Order ID" : "TH-2015-107944",
    "Order Date" : "12/12/2012",
    "Ship Date" : "22/12/2012",
    "Ship Mode" : "Standard Class",
    "Customer ID" : "TT-10210",
    "Customer Name" : "Thanaboon Tikaew",
    "Segment" : "Consumer",
    "Country" : "Thailand",
    "City" : "MaeMoh",
    "Postal Code" : 52220,
    "Region" : "North",
    "Product ID" : "OFF-MM-10000353",
    "Category" : "Furniture",
    "Sub-Category" : "Bookcases",
    "Product Name" : "Alien Cat",
    "Sales" : 23.22,
    "Quantity" : 1,
    "Discount" : 0.00,
    "Profit" : 3.11
}
my_df = pd.DataFrame([mydata])
superstore_order = pd.concat([superstore_order, my_df], ignore_index=True)
superstore_order.tail()
```

Out[]:

		Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Seg
8876	8.877000e+03	US-2016-141264	13/08/2016	19/08/2016	Standard Class	CT-11995	Carol Triggs	Cons	
8877	8.878000e+03	CA-2017-126928	17/09/2017	23/09/2017	Standard Class	GZ-14470	Gary Zandusky	Cons	
8878	8.879000e+03	CA-2017-126928	17/09/2017	23/09/2017	Standard Class	GZ-14470	Gary Zandusky	Cons	
8879	8.880000e+03	US-2015-107944	23/03/2015	25/03/2015	First Class	AM-10360	Alice McCarthy	Corp	
8880	6.707050e+10	TH-2015-107944	12/12/2012	22/12/2012	Standard Class	TT-10210	Thanaboon Tikaew	Cons	

5 rows × 22 columns



Groupby

In []: `superstore_order["Profit Ratio"] = superstore_order["Profit"]/superstore_order["Sales"]`
`superstore_order.groupby(["Category", "Sub-Category"]).agg(mean_profit_ratio = ("Pr`

Out[]:

mean_profit_ratio

Category	Sub-Category	mean_profit_ratio
Furniture	Bookcases	-0.126473
	Chairs	0.045028
	Furnishings	0.140782
	Tables	-0.147916
Office Supplies	Appliances	-0.145513
	Art	0.251678
	Binders	-0.191641
	Envelopes	0.421913
	Fasteners	0.301157
	Labels	0.429984
	Paper	0.425586
	Storage	0.092382
	Supplies	0.104970
Technology	Accessories	0.219012
	Copiers	0.317826
	Machines	-0.059535
	Phones	0.118926

[Q8] Describe an information that the result above conveys?

Ans: Mean profit ratio of each product grouping by Category and Sub-Category

Pivot and Melt

Pivot

```
In [ ]: superstore_order.pivot_table(index="State", columns="Ship Mode", values="Order ID",
```

Out[]:

	Ship Mode	First Class	Same Day	Second Class	Standard Class
State					
Alabama	9.0	1.0	18.0	30.0	
Arizona	42.0	15.0	22.0	123.0	
Arkansas	10.0	2.0	8.0	35.0	
California	302.0	106.0	346.0	1000.0	
Colorado	43.0	5.0	32.0	95.0	
Connecticut	19.0	8.0	11.0	39.0	
Delaware	16.0	2.0	13.0	55.0	
District of Columbia	0.0	0.0	3.0	7.0	
Florida	47.0	25.0	57.0	210.0	
Georgia	19.0	15.0	31.0	108.0	

In []: `pivot_table_result = superstore_order.pivot_table(index="State", columns="Ship Mode
print(pivot_table_result)`

Ship Mode	First Class	Same Day	Second Class	Standard Class
State				
Alabama	9.0	1.0	18.0	30.0
Arizona	42.0	15.0	22.0	123.0
Arkansas	10.0	2.0	8.0	35.0
California	302.0	106.0	346.0	1000.0
Colorado	43.0	5.0	32.0	95.0
Connecticut	19.0	8.0	11.0	39.0
Delaware	16.0	2.0	13.0	55.0
District of Columbia	0.0	0.0	3.0	7.0
Florida	47.0	25.0	57.0	210.0
Georgia	19.0	15.0	31.0	108.0
Idaho	3.0	0.0	2.0	13.0
Illinois	58.0	24.0	96.0	249.0
Indiana	13.0	3.0	30.0	79.0
Iowa	1.0	1.0	4.0	17.0
Kansas	6.0	1.0	2.0	15.0
Kentucky	12.0	5.0	49.0	62.0
Louisiana	7.0	2.0	14.0	15.0
Maine	0.0	0.0	0.0	5.0
Maryland	18.0	7.0	12.0	63.0
Massachusetts	14.0	4.0	35.0	71.0
Michigan	20.0	16.0	43.0	151.0
Minnesota	9.0	4.0	13.0	59.0
Mississippi	3.0	4.0	7.0	36.0
Missouri	7.0	2.0	20.0	24.0
Montana	1.0	1.0	0.0	13.0
Nebraska	6.0	3.0	6.0	20.0
Nevada	4.0	1.0	12.0	17.0
New Hampshire	2.0	0.0	10.0	13.0
New Jersey	5.0	1.0	20.0	87.0
New Mexico	1.0	0.0	9.0	22.0
New York	155.0	57.0	183.0	606.0
North Carolina	36.0	14.0	40.0	139.0
North Dakota	0.0	0.0	5.0	2.0
Ohio	66.0	47.0	84.0	199.0
Oklahoma	5.0	6.0	7.0	44.0
Oregon	20.0	0.0	15.0	81.0
Pennsylvania	103.0	9.0	78.0	341.0
Rhode Island	16.0	0.0	21.0	16.0
South Carolina	3.0	5.0	18.0	16.0
South Dakota	2.0	0.0	0.0	9.0
Tennessee	21.0	2.0	24.0	118.0
Texas	125.0	37.0	161.0	537.0
Utah	4.0	2.0	19.0	28.0
Vermont	0.0	0.0	1.0	2.0
Virginia	39.0	4.0	33.0	115.0
Washington	56.0	34.0	97.0	265.0
West Virginia	0.0	0.0	0.0	3.0
Wisconsin	12.0	3.0	10.0	66.0
Wyoming	0.0	0.0	0.0	1.0

Melt

```
In [ ]: melted_result = pd.melt(pivot_table_result.reset_index(), id_vars=["State"], var_na
```

```
print(melted_result)
```

	State	Ship Mode	Order Count
0	Alabama	First Class	9.0
1	Arizona	First Class	42.0
2	Arkansas	First Class	10.0
3	California	First Class	302.0
4	Colorado	First Class	43.0
..
191	Virginia	Standard Class	115.0
192	Washington	Standard Class	265.0
193	West Virginia	Standard Class	3.0
194	Wisconsin	Standard Class	66.0
195	Wyoming	Standard Class	1.0

[196 rows x 3 columns]

[Q9] What is the advantage of using `melt` ?

Ans : Its tidy data from wide format to long format which easier to read

[Q10] From the superstore_order, display the ascending order considering values in the 'Profit' column to group the 'Category'.

In []:

```
#enter your code
category_profit = superstore_order.groupby('Category')['Profit'].sum()
category_profit.sort_values(ascending=True)
```

Out[]:

Category	
Furniture	16861.6719
Office Supplies	105827.0238
Technology	133410.4932
Name: Profit, dtype: float64	

[Q11] Create a new column that calculates the total price (sale*quantity) before discount then group by 'product id' and 'category', then show the mean of the total price

In []:

```
#enter your code here
superstore_order["Total Price"] = superstore_order["Sales"] * superstore_order["Qua
Grouped_mean_price = superstore_order.groupby(["Product ID", "Category"])["Total Pri
Grouped_mean_price
```

```
Out[ ]: Product ID      Category
FUR-BO-10000112    Furniture    7426.566000
FUR-BO-10000330    Furniture    1258.192000
FUR-BO-10000362    Furniture    1726.898000
FUR-BO-10000468    Furniture    426.532400
FUR-BO-10000711    Furniture    3194.100000
...
TEC-PH-10004912    Technology   747.320000
TEC-PH-10004922    Technology   673.249500
TEC-PH-10004924    Technology   57.149333
TEC-PH-10004959    Technology   412.009000
TEC-PH-10004977    Technology   2441.475429
Name: Total Price, Length: 1847, dtype: float64
```

[Q12] Complete the function to apply `ratio` column that calculates from `First Class` and `Standard Class` columns on `pivot_table_result`

```
In [ ]: # function to transform the ratio
def get_class_ratio(row):

    # get the first class column
    first_class = row["First Class"]

    # get the standard class column
    standard_class = row["Standard Class"]

    # calculate the ratio
    ratio = first_class / standard_class

    return ratio

pivot_table_result["ratio"] = pivot_table_result.apply(get_class_ratio, axis=1)

pivot_table_result.head()
```

State	Ship Mode	First Class	Same Day	Second Class	Standard Class	ratio
Alabama	9.0	1.0	18.0	30.0	0.300000	
Arizona	42.0	15.0	22.0	123.0	0.341463	
Arkansas	10.0	2.0	8.0	35.0	0.285714	
California	302.0	106.0	346.0	1000.0	0.302000	
Colorado	43.0	5.0	32.0	95.0	0.452632	

[Q13] After complete Q13, What does the `apply` function do?

Ans: Apply function calls "get_class_ratio" on each row of pivot table result also "axis = 1" mean we calls by row

[Q14] Create a new column(`short_ratio`) that works the same as Q12 but with `lambda` function

```
In [ ]: pivot_table_result["short_ratio"] = pivot_table_result.apply(lambda row: row["First Class"] / row["Second Class"], axis=1)
pivot_table_result.head()
```

```
Out[ ]: Ship Mode  First Class  Same Day  Second Class  Standard Class    ratio  short_ratio
State
Alabama          9.0        1.0       18.0        30.0  0.300000  0.300000
Arizona          42.0       15.0       22.0       123.0  0.341463  0.341463
Arkansas         10.0        2.0        8.0        35.0  0.285714  0.285714
California      302.0      106.0      346.0      1000.0  0.302000  0.302000
Colorado          43.0        5.0       32.0        95.0  0.452632  0.452632
```

[Q15] What is the difference between `apply` and `lambda` function? give 2 examples use case.

Ans : "apply" is Pandas method that operate function in each row of Dataframe while "lambda" is Python function that uses on Math/Logic calculation

Ex1: In `superstore_order["Customer Name"].apply(len)` we called "len" function to find length of Customer name in each row

```
In [ ]: superstore_order["Customer Name"].apply(len)
```

```
Out[ ]: 0      11
1      11
2      15
3      13
4      13
..
8876   12
8877   13
8878   13
8879   14
8880   16
Name: Customer Name, Length: 8881, dtype: int64
```

Ex2: In `superstore_order.apply(lambda row: row["Total Price"]/100, axis=1)`
we use "apply" to access all row of `superstore_order`
and in each row we use lambda to calculate its Totalprice/100 without using define function

```
In [ ]: superstore_order.apply(lambda row: row["Total Price"]/100, axis=1)
```

```
Out[ ]: 0      5.239200
        1      21.958200
        2      0.292400
        3      47.878875
        4      0.447360
        ...
8876    0.589240
8877    19.200000
8878    1.021500
8879    21.199200
8880    0.232200
Length: 8881, dtype: float64
```